

SE 2XB3 Group 4 Report 3

Huang, Kehao	Jiao, Anhao
400235182	400251837
huangk53@mcmaster.ca	jiaoa3@mcmaster.ca
L01	L01

Ye, Xunzhou
400268576
yex33@mcmaster.ca
L01

5 February 2021

1 Quicksort

1.1 In-Place Version

Our implementation has its natural in-place advantage over the given implementation which uses auxiliary memory to store different partitions for each recursion call. Specifically, our in-place implementation swaps elements in the array and uses two parameter `low` and `high` to mark the array partition on which a recursion call should work. The amount of memory used for the whole sorting process is independent of the input size. In this case, it is the length of the input array. On the other hand, the given non-in-place implementation copies elements from the input array to fresh allocated auxiliary arrays. Each auxiliary array is used as the new partition for the subsequent recursion call. And the returned sorted array is a concatenation of two sorted partitions and the pivot. Both the element copying action and list concatenation are costly, in terms of both time and space complexity.

A test on the average runtime of both versions of quicksort is then carried out. For n on the scale from 10^4 to 10^7 , an array of n random numbers is passed to both implementations. The runtime is plotted on a semi-log graph as shown in Figure 1. The experimental result did not precisely match our prediction. From $n = 10^4$ to approximately $n = 10^{5.5}$, the in-place version has a slight longer runtime than the non-in-place version. Figure 2 is a zoomed-in view of the plot, which demonstrates the shorter runtime of the non-in-place version. As n increases from $10^{5.5}$, the in-place quicksort gradually starts to outperform the other. The disadvantage of using auxiliary arrays and deep copying data around memory becomes obvious as n passes 10^6 .

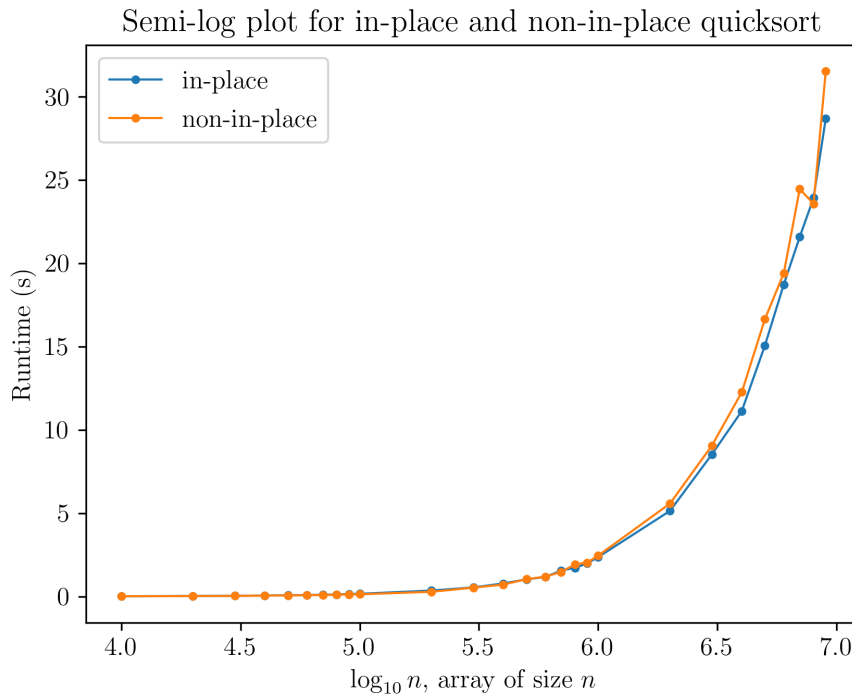


Figure 1: In-place and non-in-place quicksort comparison




Figure 2: Quicksort versions comparison on lower scale n

To quantify the performance difference between the two versions, we used a rather simple model. Either version of quicksort is known to have a complexity of $\mathcal{O}(n \lg n)$. We use $c * n \lg n$ as an approximate model of the imperial runtime. Thus we take the average of the differences of the c constants of the two versions as the quantified performance difference. From the same data set plotted above, for $10^4 \leq n < 10^7$, the non-in-place version is on average 4% faster than the in-place version.

In conclusion, the non-in-place quicksort is slightly faster than the in-place one in practice. However, the in-place version has an observable speed improvement for an input size $n > 10^6$.

1.2 Multi-Pivot

Variants of quicksorts with different numbers of pivots were implemented and tested against the traditional single-pivot quicksort. It is worth noting that within the implementations of the quicksorts with two or more pivots, the provided single-pivot quicksort is used to sort the chosen pivots and any input array of length less than the number of pivots.

The performance of the quicksort variants was tested using arrays of length $10^4 \leq n < 10^7$ (input size). As a result, the quad-pivot quicksort outperformed the other sorting algorithms and was chosen as the recommended quicksort. As shown in Figure 3, given an input size, the quicksort finished faster as the number of pivots increased. The runtime differences between variants widened as the input size increased. However, the time advantage of the quad-pivot variant is only significant in the input range $n > 10^6$. Figure 4 provide the zoomed-in views for $10^4 \leq n \leq 10^5$ and $10^4 \leq n \leq 10^6$. The performance of all algorithms is rather unstable in these two ranges. There is no clear superior out of the four tested variants. Therefore, the

quad-pivot quicksort is concluded to be the fastest solely based on its better performance on large sized inputs.

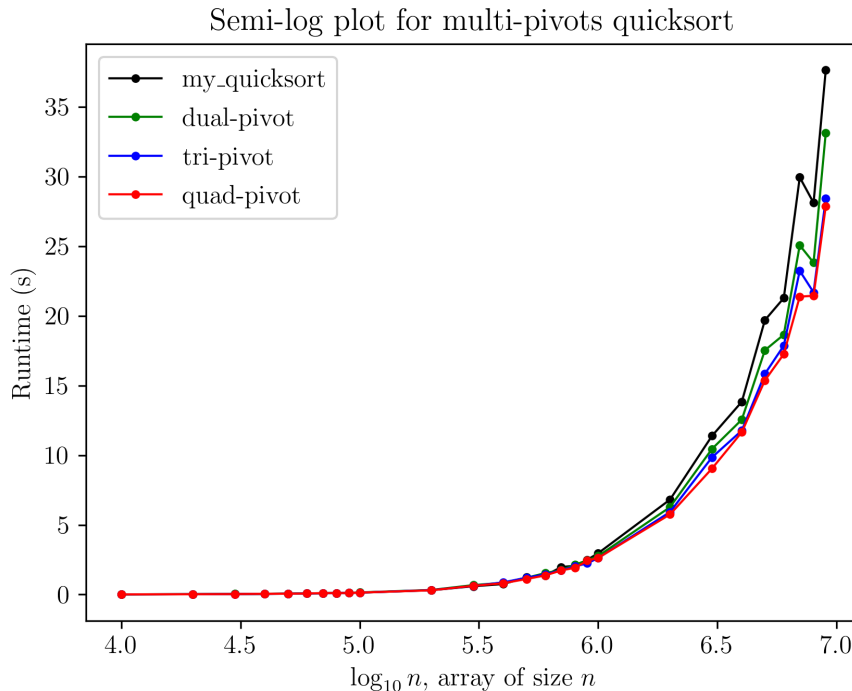


Figure 3: Multi-pivot quicksorts runtimes

1.3 Worstcase Performance

The worstcase performance of quicksort happens when the input list is already sorted in either ascending or descending order and when the implementation always chooses the first or the last element as the pivot. In this case, the algorithm would only be able to split the list down to two extremely unbalanced partitions: one with zero length and the other with length $n - 1$. This partitioning scheme results in a $\sim n$ recursion depth. Also, in each recursion call, $\sim n$ comparisons are required to complete the partitioning step. The overall worstcase performance is expected to be $\mathcal{O}(n^2)$. The timing experiment result shown in Figure 5 demonstrates the significant difference between quicksort on average case and worstcase.

For near-sorted lists of elements, insertion sort is expected to outperform quicksort. This is because quicksort has extra overhead from the recursive calls, while the core of the insertion sort is simply a **for** loop. Focusing on lists of length 1000, an experiment was conducted to compare the runtime of four types of sorting algorithms by varying the inversion factor of an input list. The inversion factor is a percentage of how many elements in a list are inverted in order. As discussed previously, a sorted (zero inversion) list is the worstcase for quicksort. Figure 6 reflects the relatively slow runtime of quicksort for lower inversion factors. However,

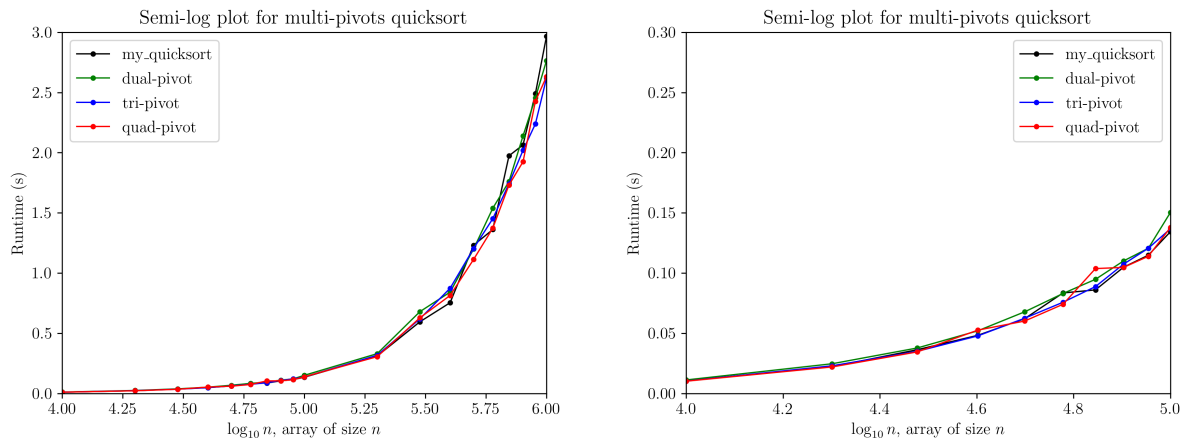


Figure 4: Quicksort variants on lower scale n

quicksort took the lead at around 5% inversion factor and performed significantly better than other elementary sorting algorithms for higher inversion factors.

1.4 Small Lists

The experiment focused on small size lists, for input size $1 \leq n \leq 20$, the insertion sort performed the best compared to other three sorting algorithms as shown in Figure 7. A small list has fewer unsorted elements than a large list, which shares some similarity with a near sorted list. And the insertion sort performed similarly in both two experiments. Therefore, insertion sort is the optimal approach while handling small size lists.

To summarize all the experiments above, we decided to define our `final_sort` to be a hybrid of insertion sort and quad-pivot quicksort. When the length of the input list is sufficiently short (less than 20), we would use insertion sort. For any other lists with longer lengths, we would use quad-pivot quicksort.

The purpose of designing this hybrid algorithm is to minimize the average runtime. From Section 1.1, even though the in-place version of quicksort is faster than the non-in-place version on large scaled lists, it is still slower than the quad-pivot variant tested in Section 1.2. Furthermore, from Section 1.4, insertion sort has an observable advantage over other algorithms in terms of runtime for lists of length 20 and below. As for the lessons learned from Section 1.3, near-sorted lists in descending order are worstcase for both insertion sort and quicksort with $\mathcal{O}(n^2)$ complexity, while ascending ordered lists are best case for insertion sort but worstcase for quicksort. An ideal insertion-quicksort hybrid sorting algorithm would use insertion sort's best case to balance out quicksort's worstcase for ascending sorted lists. However, there is not sufficiently efficient algorithm to determine whether an input list is near-sorted. One possible solution to balance the worstcase for quicksort is to use a randomized quicksort. However, such implementation is beyond the scope of the lab. Therefore, we conclude that a hybrid sorting algorithm comprises insertion sort and quicksort which conditionally chooses a sorting algorithm based on the input size, is the best performing sorting algorithm on average case.




Figure 5: Quicksort on average case and worstcase

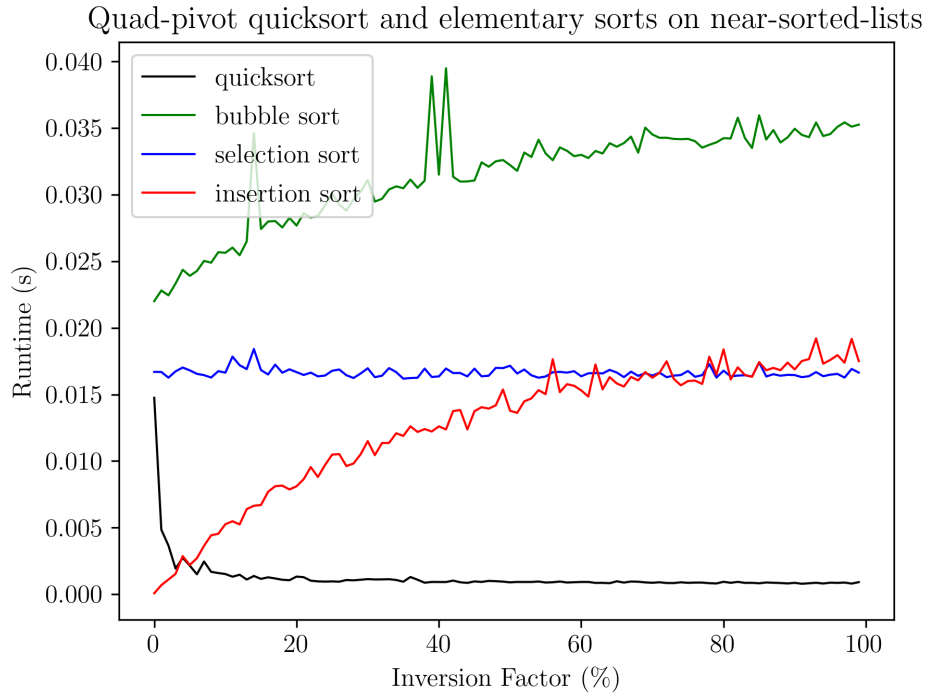


Figure 6: Runtime of different sorting algorithms on near-sorted-lists

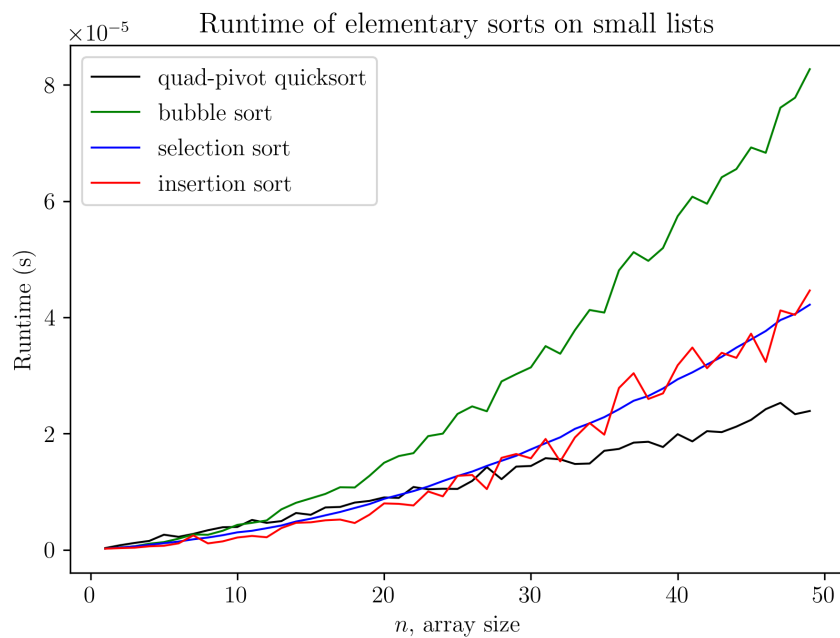


Figure 7: Runtime of different sorting algorithms on small lists